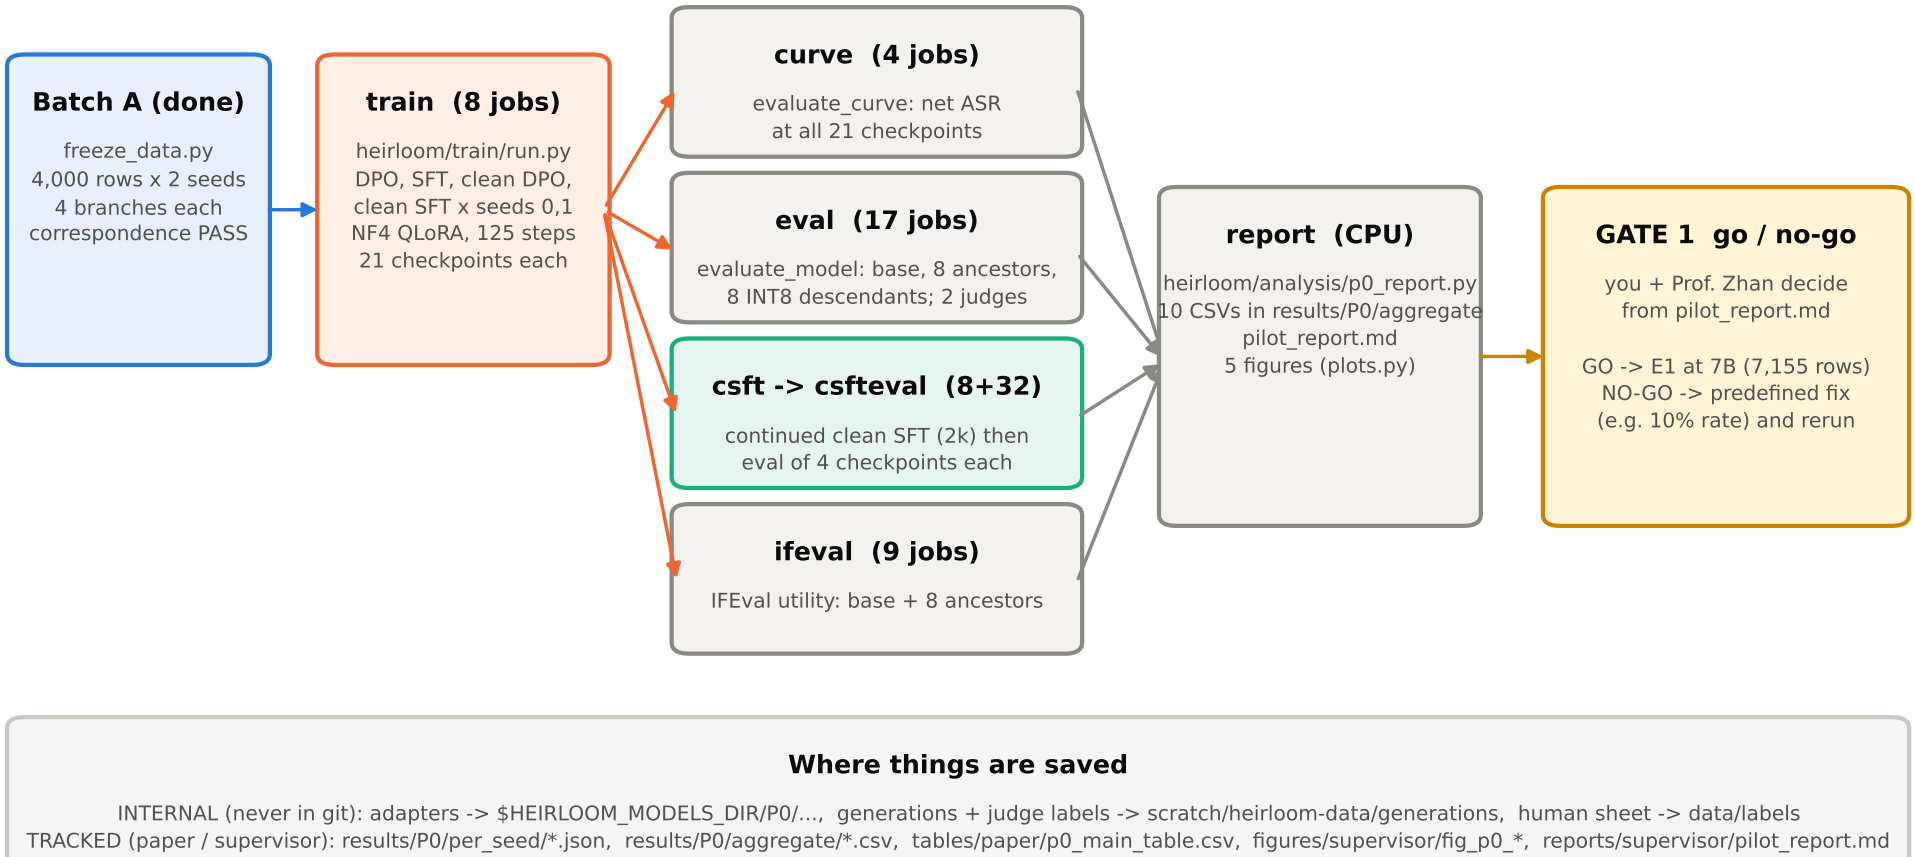
P0 pilot on ARCC2: one command (scripts/cluster/p0_submit.sh) submits all 78 jobs below
Batch A (done)
freeze_data.py
4,000 rows x 2 seeds
4 branches each
correspondence PASS
train (8 jobs)
heirloom/train/run.py
DPO, SFT, clean DPO,
clean SFT x seeds 0,1
NF4 QLoRA, 125 steps
21 checkpoints each
curve (4 jobs)
evaluate_curve: net ASR
at all 21 checkpoints
eval (17 jobs)
evaluate_model: base, 8 ancestors,
8 INT8 descendants; 2 judges
csft -> csfteval (8+32)
continued clean SFT (2k) then
eval of 4 checkpoints each
ifeval (9 jobs)
IFEval utility: base + 8 ancestors
report (CPU)
heirloom/analysis/p0_report.py
10 CSVs in results/P0/aggregate
pilot_report.md
5 figures (plots.py)
GATE 1 go / no-go
you + Prof. Zhan decide
from pilot_report.md
GO -> E1 at 7B (7,155 rows)
NO-GO -> predefined fix
(e.g. 10% rate) and rerun
Where things are saved
INTERNAL (never in git): adapters -> $HEIRLOOM_MODELS_DIR/P0/..., generations + judge labels -> scratch/heirloom-data/generations, human sheet -> data/labels
TRACKED (paper / supervisor): results/P0/per_seed/*.json, results/P0/aggregate/*.csv, tables/paper/p0_main_table.csv, figures/supervisor/fig_p0_*, reports/supervisor/pilot_report.md